

Edge Detection

CS 450: Computer Vision
slides by Dr. Bryan Morse

Edges and Gradients

- ▶ Edge: *local* indication of an object transition
- ▶ Edge detection: local operators that “find” edges (usually involves convolution)
- ▶ Local intensity transitions are indicated by the *gradient*:

$$\nabla I = \begin{bmatrix} \frac{\partial}{\partial x} I \\ \frac{\partial}{\partial y} I \end{bmatrix}$$

- ▶ Interpretation:
 - ▶ Gradient magnitude $\|\nabla I\|$: edge “strength”
 - ▶ Gradient orientation $\phi(\nabla I)$: cross-edge direction

Prewitt Kernels

Idea: central finite differences

Central difference:

$$\frac{dl}{dx} \approx [I(x+1) - I(x-1)]/2$$

or for two-dimensional images:

$$\frac{\partial I}{\partial x} \approx [I(x+1, y) - I(x-1, y)]/2$$

This corresponds to the following convolution kernel:

-1	0	1	-1
			0
			1

Prewitt Kernels

Or for more robustness to noise, smooth in the other direction:

-1	0	1
-1	0	1
-1	0	1

 $\partial/\partial x$

-1	-1	-1
0	0	0
1	1	1

 $\partial/\partial y$

Sobel Kernels

Or, giving more weight to the central pixels when averaging:

-1	0	1
-2	0	2
-1	0	1

 $\partial/\partial x$

-1	-2	-1
0	0	0
1	2	1

 $\partial/\partial y$

- ▶ These kernels can also be thought of as 3×3 approximations of the first derivative of a small Gaussian
- ▶ Can be thought of as blurring by a small Gaussian to remove noise, then taking the derivative:

$$\frac{\partial}{\partial x}(I * G) = I * \frac{\partial}{\partial x} G \approx I * \text{Sobel}(x)$$

From Gradient Magnitude to Edges

The gradient magnitude gives a measure *at every pixel* of the “edginess” of each pixel:

$$\|\nabla I(x, y)\|$$

Somehow, you have to still find the *best* edges:

- ▶ Threshold (global or local), etc.
- ▶ Local maxima of gradient magnitude

⋮

Using Second Derivatives

- ▶ Classic maximum test from calculus:

$$x \text{ is a extremum of } f(x) \text{ if } \frac{df}{dx}(x) = 0$$

- ▶ Extend this idea to find maximal first derivatives:




$$x \text{ is a extremum of } \frac{df}{dx}(x) \text{ if } \frac{df^2}{dx^2}(x) = 0$$

Using Second Derivatives

a b
c

FIGURE 3.38

(a) A simple image, (b) 1-D horizontal gray-level profile along the center of the image and including the isolated noise point, (c) Simplified profile (the points are joined by dashed lines to simplify interpretation).



Laplacian

- ▶ The *Laplacian* is defined mathematically as

$$\nabla^2 = \nabla \cdot \nabla = \begin{bmatrix} \frac{\partial}{\partial x} \\ \frac{\partial}{\partial y} \end{bmatrix} \cdot \begin{bmatrix} \frac{\partial}{\partial x} \\ \frac{\partial}{\partial y} \end{bmatrix} = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$$

- ▶ When we apply it to an image, we get

$$\nabla^2 I = \left(\begin{bmatrix} \frac{\partial}{\partial x} \\ \frac{\partial}{\partial y} \end{bmatrix} \cdot \begin{bmatrix} \frac{\partial}{\partial x} \\ \frac{\partial}{\partial y} \end{bmatrix} \right) I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$

Marr-Hildreth Edge Detection

Idea: approximate finding maxima of gradient magnitude (edges) by finding places where

$$\nabla^2 I(x, y) = 0$$

- ▶ Can't always find discrete pixels where the Laplacian is *exactly* zero—look for *zero crossings* instead.

Laplacian Operators

- ▶ Second difference:

$$\begin{aligned}\frac{d^2 I}{dx^2} &\approx [I(x+1) - I(x)] - [I(x) - I(x-1)] \\ &= I(x+1) - 2I(x) + I(x-1)\end{aligned}$$

- ▶ The Laplacian is one of these in the x direction added to one of these in the y direction:

0	0	0
1	-2	1
0	0	0

+

0	1	0
0	-2	0
0	1	0

=

0	1	0
1	-4	1
0	1	0

Derivatives of Gaussians

- ▶ Gaussian kernel for noise removal:

$$G(x, y) = (2\pi\sigma)^{-d/2} e^{-r^2/2\sigma^2}$$

where $r^2 = x^2 + y^2$ and d is the dimensionality (images=2)

- ▶ Can solve in closed form for the first- and second-order derivatives of the Gaussian (including the Laplacian)
- ▶ Can convolve with these directly—
exactly the same as blurring first then applying derivatives

Difference of Gaussians

- ▶ Another property of the Laplacian of Gaussian:

$$\nabla^2 G = \frac{\partial}{\partial \sigma} G$$

- ▶ We can thus approximate the Laplacian by the difference of one Gaussian and a just-smaller one:

$$\nabla^2 G \approx G(x, y; \sigma_1) - G(x, y; \sigma_2)$$

This is the *Difference of Gaussians* (DoG) kernel.

- ▶ Ratio (σ_1/σ_2) for best approximation is about 1.6.
(Some people like $\sqrt{2}$.)

Combining Both First- and Second-Derivatives

- ▶ Laplacian zero crossings:

$$\nabla^2 I = 0$$

- ▶ Problem:

Tells you the gradient magnitude is at a maximum, *not* how strong it is—lots of spurious edges.

- ▶ Idea:

Combine the two measures

$$\nabla^2 I = 0 \text{ and } \nabla I > T$$

Canny Edge Detector

- ▶ Problem with Laplacian zero-crossings: adds the principal curvatures together—it doesn't really determine a maximum of gradient magnitude in any one direction.
- ▶ The *Canny Edge Detector* defines edges as *zero-crossings of second derivatives in the direction of greatest first derivative*.

The Direction of Greatest First Derivative

This is simply the gradient direction.

The Second Derivative in The Direction of . . .

We can compute this using the matrix of second derivatives (Hessian).

Zero Crossings of . . .

As with the Marr-Hildreth edge detector, we'll use positive-negative transitions to "trap" zeroes.

- ▶ Gives connected edges much like the Laplacian operator but more accurately localize the edge.

Now What?

Now you have potential edge points, how do you get contours?

- ▶ Thresholding gradient magnitude
- ▶ Threshold, then “relax” the labeling
 - ▶ Thresholding with hysteresis (Canny)
 - ▶ Edge relaxation algorithms using other criteria
- ▶ Edge linking (including postprocessing)
- ▶ Connected loci of local maxima (ridges)
- ▶ Maximum-magnitude contours/paths
(Turns into optimization problem—we'll come back later.)

Representation

Once you have contours, how do you represent them?

- ▶ Chain codes (4- or 8-connected directions)
- ▶ Differential chain codes (4- or 8-connected *relative* directions)
- ▶ Polylines (fit with short line segments)
- ▶ Arc-length parameterization
 - ▶ Position
 - ▶ Tangent orientation
 - ▶ Curvature
 - ▶ Distance to some central point
 - ▶ ...
- ▶ Fourier descriptors
- ▶ Many other ways to represent curves

Representation

What do you want to do with the representation?

- ▶ Reproduction
- ▶ Matching
 - ▶ Translated
 - ▶ Rotated
 - ▶ Scaled
 - ▶ Noise and variation (smooth the curves)
 - ▶ ...

Want to be *invariant* to as much as possible